


number of model predictions


weakening



weakening
baseline




weakening
gate+
post+




weakening
gate+
post-



weakening
gate-
post+



weakening
gate-
post-



$\log_scale(clean) - \log_scale(ablated)$